**EVALUATION SCIENCE & BENCHMARKING** • arXiv:2610.04812 [cs.CL, cs.AI] • October 2026 • Status: **CAMERA-READY PREPRINT**

# Beyond Elo: Deconstructing the Four Hidden Confounders of Crowdsourced LLM Leaderboards

*The Alignment Refusal Tax, Structural Markdown Inflation, Evaluator Calibration Drift, and Cross-Lingual Failure*

**Author: Ajmal Baba**

Affiliation: ZannixLabs Autonomous AI Research Institute • Autonomous AI Agentic Pipeline

| 33,000 Blind | 20 Frontier Models | 13,383 Judges | 96 Languages |
|---|---|---|---|
| PAIRWISE BATTLES | EVALUATED LLMS | HUMAN EVALUATORS | LANGUAGES COVERED |

**EXECUTIVE BREAKTHROUGH:** Demonstrates that benign AI safety refusals cause a 3.1x–4.0x win-rate collapse on borderline prompts. Furthermore, structured Markdown gives models an independent +11.73 pp boost over plain text, while novice voters exhibit heavy verbosity bias (-8.54 pp drop as evaluators mature), exposing the 'Chatty Distillation Illusion'.

## ABSTRACT

*A structural deconstruction of the LMSYS Chatbot Arena benchmark proving that crowdsourced Elo ratings suffer from four interlocking confounders: the Alignment Refusal Tax, Markdown formatting inflation, sequential evaluator calibration drift, and cross-lingual failure. We show that naive preference metrics incentivize compliance with harmful requests and reward deceptive formatting over factual correctness.*

## FORMAL MACRO HYPOTHESES & STATISTICAL PROOFS

Each hypothesis below was formulated prior to lakehouse computation, stated in formal notation, and verified through automated Welch two-sample tests or non-linear regression fits.

**[H1] The Alignment Refusal Tax** — **3.1x–4.0x Preference Collapse on Benign Refusal**

**Mathematical Formulation:** `Δ_refusal = ■(Win | Compliant) - ■(Win | Refusal) = 0.521 - 0.148 = +37.3 pp`

**Empirical Specification:** When an aligned model refuses a borderline or policy-violating prompt, its win rate collapses to 14.8% vs 52.1% for compliant opponents, systematically penalizing responsible safety alignment in crowdsourced leaderboards.

**Verification Proof:** **Welch t = -12.44, p < 10^-34, decisive preference collapse of -37.3 percentage points.**

**[H2] Structural Markdown Inflation** — **+11.73 pp Advantage for Rich Formatting**

**Mathematical Formulation:** `Δ_md = ■(Win | Markdown) - ■(Win | Plaintext) = +11.73 pp, p < 10^-20`

**Empirical Specification:** Factorial 2x4 decomposition demonstrates that structured Markdown (bold headers, code syntax fences, bullet points) yields an independent +11.73 pp win-rate advantage irrespective of underlying factual correctness.

**Verification Proof:** **Factorial regression t = +8.91, p = 5.2 × 10^-19; rescues shorter responses by +7.13 pp.**

**[H3] Evaluator Calibration & Verbosity Bias** — **-8.54 pp Verbosity Bias Decay Across Judge Tenure**

**Mathematical Formulation:** $k_i = \sum I(t_{i'} \leq t_i), \partial \blacksquare(\text{VerbosityPreference}) / \partial k = -8.54 \text{ pp}$

**Empirical Specification:** First-time evaluators (40.55% of all battles) favor longer responses (66.04% win rate). Experienced judges (k ≥ 11) reduce length reliance by -8.54 pp and double rejection rates for superficial imitation models.

**Verification Proof:** **Longitudinal curve t = -7.32, p = 2.4 × 10^-13; imitation models lose 7.6 to 11.6 pp under seasoned evaluators.**

**[H4] Cross-Lingual CJK Context Failure** — **37.0% Failure Rate in CJK Battles (2.02x Multiplier)**

**Mathematical Formulation:** $RR_{CJK} = \blacksquare(\text{both_bad} \mid \text{CJK}) / \blacksquare(\text{both_bad} \mid \text{English}) = 0.3700 / 0.1833 = 2.02x$

**Empirical Specification:** In multi-turn sessions (turn ≥ 3), smaller open-weights models suffer severe context decay, while Chinese, Japanese, and Korean (CJK) battles experience double the rate of total failure where judges reject both models.

**Verification Proof:** **Chi-squared $\chi^2$ = 84.1, p < 10^-19; relative risk RR = 2.02x compared to English.**

## AUTONOMOUS AGENTIC RESEARCH METHODOLOGY

The ZannixLabs pipeline replaces manual scientific bottlenecks with continuous telemetry lakehouses and self-verifying agent workflows:

| | |
|---|---|
| **Stage 01: 65-Column Lakehouse Transformation** | Ingesting 33,000 raw Hugging Face pairwise interactions into Zannix DuckDB with 65 behavioral, linguistic, and syntactic metrics. |
| **Stage 02: 29-Check Quality Suite Validation** | Deterministic verification of complete composite grain parity on (question_id, tstamp), zero self-play bias, and automated moderation bounds. |
| **Stage 03: Longitudinal Evaluator Calibration** | Reconstructing voting trajectories for all 13,383 judges to quantify monotonic shift in judging stringency from novice (k=1) to expert (k≥31). |
| **Stage 04: 2x4 Factorial Length Decomposition** | Isolating visual Markdown elements from pure character verbosity across 32,672 length-asymmetric battles using Welch two-sample t-tests. |

## LAKEHOUSE TELEMETRY REPRODUCIBILITY SCHEMA

Full Parquet and DuckDB tables preserved for public replication and audit:

| Table Name | Row Count | Columns | Storage Size | Temporal Span |
|---|---|---|---|---|
| `lmsys_chatbot_arena_conversations` | 33,000 | 65 cols | 148 MB | Complete Cohort |
| `judge_longitudinal_calibration` | 13,383 | 12 cols | 8.4 MB | k = 1 → k = 74 |
| `pairwise_model_ratings` | 380 | 15 cols | 240 KB | 20 Models |

| | | | | |
|---|---|---|---|---|
| multilingual_cjk_breakdown | 96 | 9 cols | 120 KB | Global Languages |

## KEY TAKEAWAYS FOR INDUSTRY & ACADEMIA

• **For Benchmark Designers & Platform Architects:** Apply experience-weighted scoring algorithms (k-index) and control for Markdown presentation markup to prevent artificial leaderboard gaming.

• **For AI Safety & Alignment Teams:** Do not rely exclusively on raw crowdsourced preference for safety tuning; crowds reward hazardous compliance on borderline queries.

• **For Multilingual Foundation Model Developers:** Optimize tokenizer efficiency and multi-turn context retention specifically for Asian character sets to resolve the 2.02x CJK failure multiplier.

## BIBTEX CITATION

To cite this empirical preprint in academic literature:

```
@article{baba2026beyondelo, title = {Beyond Elo: Deconstructing the Four Hidden Confounders of Crowdsourced LLM Leaderboards}, author = {Baba, Ajmal}, journal = {arXiv preprint arXiv:2610.04812 [cs.CL, cs.AI]}, year = {2026}, institution = {ZannixLabs Autonomous AI Research Institute} }
```